

Capstone Project Reinforcement Learning: Tetris

Andrew Falcone
10/17/2022





Project Summary:

Program Tetris and program an agent to play the game.

- 1) Create the game Tetris playable for humans (N64 color schema)
- 2) Create an agent able to play the game optimally (in progress)
 - Limitations: agent is only able to choose to swap, rotate and then position piece. No rotations or slides upon landing to reduce training complexity.





Motivation

- Love of games and desire to practice implementing a reinforcement learning algorithm!
- DeepMind's recent breakthroughs using reinforcement learning algorithms including:
 - AlphaGo 2016
 - AlphaZero 2017
 - AlphaFold 2020



Meet the Pieces & Rotations





Goal of Playing the Game

1) CLEAR LINES

2) More Lines, the better (demo)

3) 4 Lines at once is a Tetris and worth more points (depends on version actual score)




Coding Challenges in Game Creation


- 1) Started with starter code game to learn about Pygame clocks and rendering
(<https://levelup.gitconnected.com/writing-tetris-in-python-2a16bddb5318>)
- 2) Fixed color scheme from random to N64
- 3) Implemented SRS rotation system
- 4) Added shadow block for piece placement
- 5) Added piece queue to see in coming pieces and the ability to swap piece
- 6) Added debugging features such as choosing which block to come next and adding lines/levels for free




SRS: Super Rotational System




1. Initial position.
Attempt to rotate $0 \gg 3$.



2. Test 1, basic rotation fails.




3. Test 2, (1, 0) fails.



4. Test 3, (1, 1) fails.



5. Test 4, (0, -2) fails.



6. Final position.
Test 5, (1, -2) succeeds.



Shadow Piece

Time: 77 s **Score: 0**

Human: 'p' to swap

Lines: 0

Level: 0

SWAP!

Controls

/: CCW rotation

rShift: CW rotation

up,down,left,right: movement


space: hard drop

0-6: debug blocks

s: swap

q: restart game

l: free line




Queue:



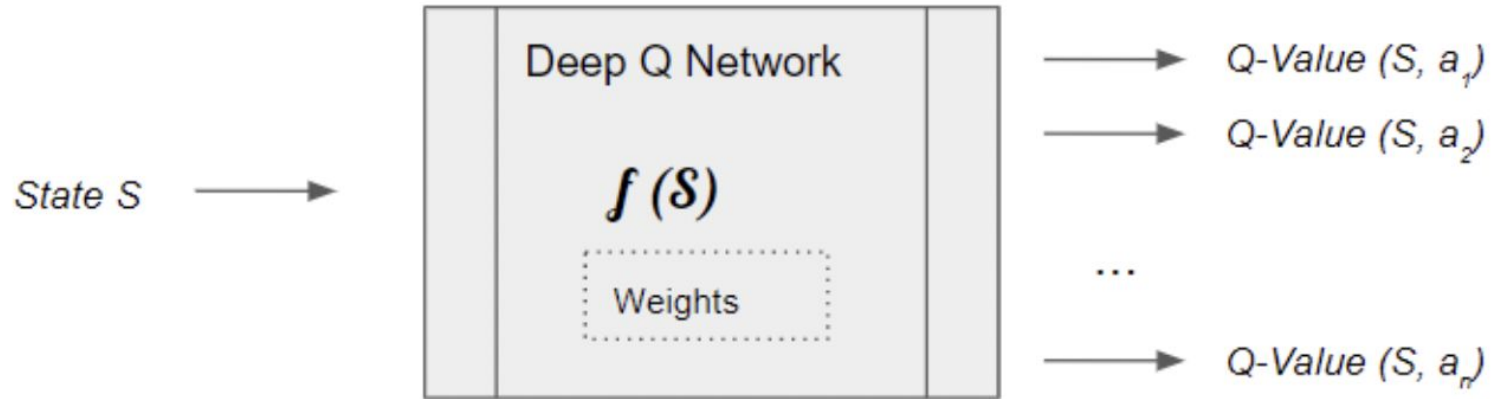


Reinforcement Learning





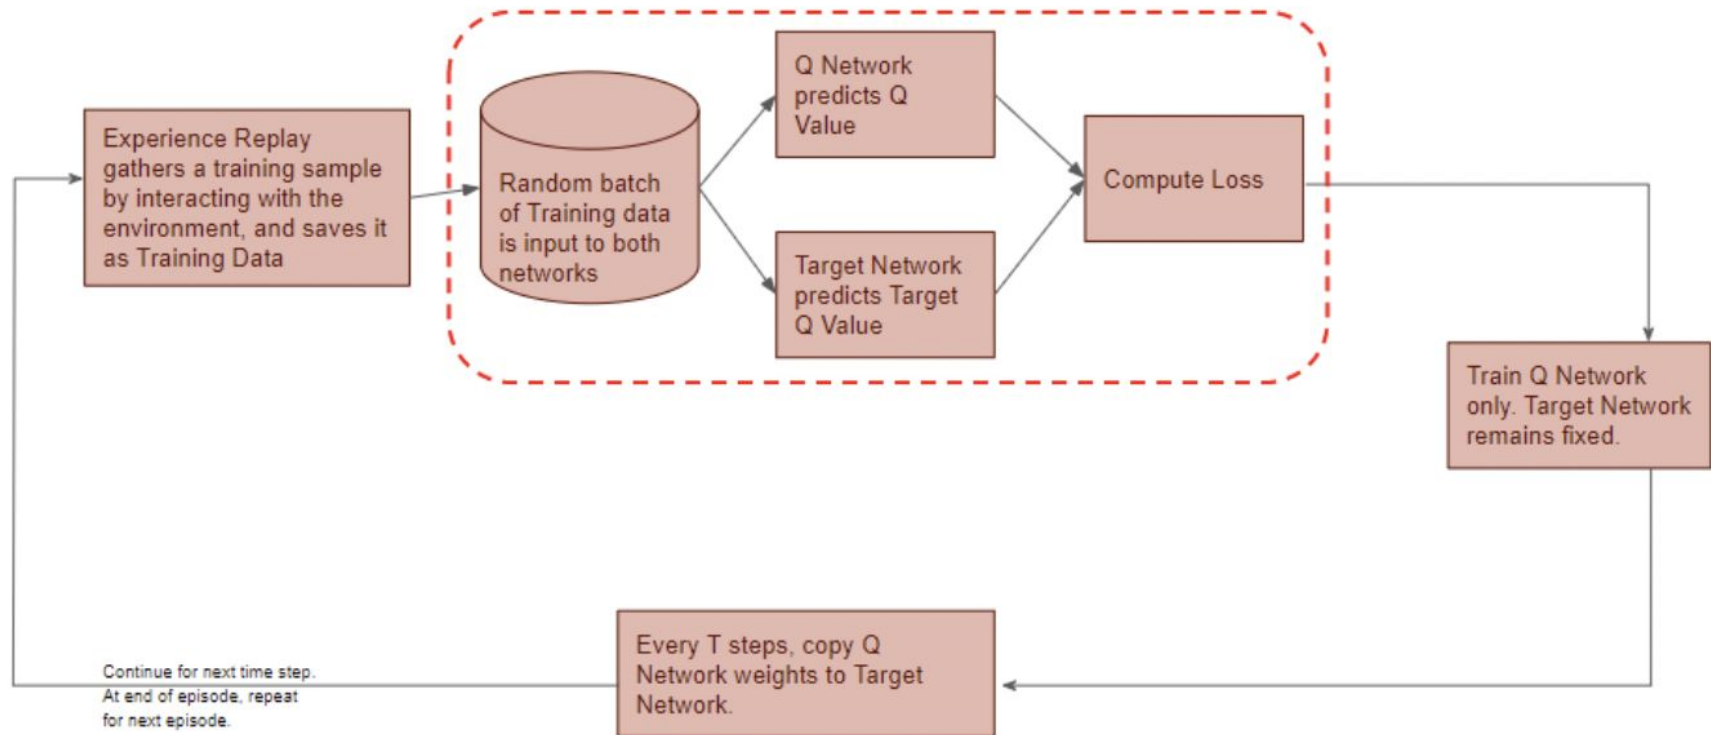


Q Network




Value Functions and Q Functions







Challenges



- 1) Working with OpenAI Gym- version stability and kernel crashes
- 2) Bugs mid-game are difficult to find and fix! Simple indexing issues are a nightmare
- 3) Exploitation vs Exploration
 - a) Agent has difficulty learning good states
- 4) Training time (agent is still training)
- 5) Getting all the pieces of code to literally fit together
 - a) Separating human vs computer play
 - b) Getting pertinent data into neural network
 - c) Implementing code in OOP style
- 6) Using too much data (whole game matrix initially), did not show any signs of converging or learning



Future Goals

Finish modeling and I hope to share a running model soon!

Learn more reinforcement learning and hopefully add extensions.

Questions???